

Searching for Axion-Like Particles at FCC-ee: Projected Sensitivity in the Invisible and Prompt-Resolved Channels at the Z Pole

Iñaki Pardo

Maastricht Science Programme, Maastricht University

inaxpardo0204@gmail.com

June 2026 — Project Period Draft

Abstract

We present projected sensitivity estimates for the production of axion-like particles (ALPs) at the proposed Future Circular Collider operating at the Z pole ($\sqrt{s} = 91.2 \text{ GeV}$, $\mathcal{L} = 150 \text{ ab}^{-1}$) via the associated production process $e^+e^- \rightarrow \gamma a$. Two detector signatures are analysed: the *invisible* channel, in which the ALP escapes the detector ($\ell_a > L_{\text{max}}$), and the *prompt-resolved* channel, in which the ALP decays to two separated photons inside the tracker ($\ell_a < L_{\text{min}}$, $\Delta\theta > \Delta\theta_{\text{res}}$). Signal events are generated with MADGRAPH5_AMC@NLO, processed through PYTHIA 8 for ALP decay and lifetime, and passed through a DELPHES fast simulation of the IDEA detector concept. SM backgrounds for both channels are simulated, binned, and normalised to the full target luminosity. Projected 90% CL exclusion contours are derived using a binned Asimov $\Delta\chi^2$ method, with analytic production and lifetime formulae corrected by detector-level signal scans. For the invisible channel, the analysis reaches $g_{a\gamma\gamma} \approx 5.5 \times 10^{-7} \text{ GeV}^{-1}$ over the ALP mass range 0.01–0.92 GeV. The prompt-resolved channel extends sensitivity to $g_{a\gamma\gamma} \sim 10^{-5} - 3 \times 10^{-4} \text{ GeV}^{-1}$ for $m_a \in [0.61, 80] \text{ GeV}$. Both projected contours surpass current collider bounds by one to two orders of magnitude in coupling at comparable masses. The analysis machinery is validated against the published Belle II ALP exclusion boundary, reproducing it to within $|\log_{10}(g_{\text{closure}}/g_{\text{published}})| \leq 7.6 \times 10^{-3}$. Additional validation gates check the MadGraph cross section, the ALP width and Pythia lifetime convention, the Delphes output, and the reconstructed ALP invariant-mass or recoil-energy observable at every detector-level signal point.

1 Introduction

The Peccei–Quinn (PQ) mechanism, proposed to explain the absence of strong CP violation [1–3], predicts the existence of a pseudo-Nambu–Goldstone boson, the axion. The simplest requirement for the PQ mechanism to resolve the strong CP problem constrains the axion mass and its couplings to a narrow band in the coupling–mass plane—the so-called QCD axion line. However, extensions of this framework that decouple the mass from the coupling give rise to *axion-like particles* (ALPs): pseudo-scalars that couple to Standard Model (SM) particles with arbitrary mass–coupling combinations [4–6].

The ALP–photon coupling,

$$\mathcal{L} \supset -\frac{1}{4}g_{a\gamma\gamma}aF_{\mu\nu}\tilde{F}^{\mu\nu}, \quad (1)$$

is particularly well-motivated. It arises generically from the anomalous coupling of any axion-like field to the electromagnetic field strength $F_{\mu\nu}$, and is bounded across many orders of magnitude in mass by astrophysical, cosmological, and laboratory experiments [12]. In the mass range 10 MeV–80 GeV relevant to collider searches, the most stringent direct constraints come from e^+e^- collider experiments (LEP, Belle II, BaBar, BESIII) and beam-dump experiments, and leave significant regions of parameter space unexplored at couplings $g_{a\gamma\gamma} \lesssim 10^{-4} \text{ GeV}^{-1}$.

The Future Circular Collider (FCC-ee) is a proposed e^+e^- machine that would operate at several centre-of-mass energies from the Z pole ($\sqrt{s} \approx 91.2 \text{ GeV}$) up to the $t\bar{t}$ threshold [8, 9]. At the Z pole, the baseline luminosity target of 150 ab^{-1} —corresponding to approximately 5×10^{12} Z bosons—provides an extraordinarily clean and prolific environment for precision measurements and rare-process searches. For photophilic ALPs, the associated production cross section $\sigma(e^+e^- \rightarrow \gamma a) \propto g_{a\gamma\gamma}^2(1 - m_a^2/s)^3$ grows quadratically with the coupling and benefits enormously from the high luminosity, making FCC-ee a uniquely powerful probe. The Z-pole run is chosen as the baseline deliverable for three practical and physics reasons. First, it is the FCC-ee stage with the largest expected integrated luminosity, so it gives the strongest statistics-driven reach for the single-coupling process studied here. Second, the fixed centre-of-mass energy gives a simple mono-energetic recoil photon, $E_\gamma = (s - m_a^2)/(2\sqrt{s})$, making the invisible and prompt-resolved analyses especially transparent. Third, the Z-pole overlaps the mass range where existing collider constraints from Belle II, BaBar, BESIII, LEP, and the LHC leave open regions that can be tested in a clean e^+e^- environment. Other FCC-ee stages are scientifically interesting, but each stage has a different luminosity, background composition, recoil-energy range, and detector acceptance optimisation. They are therefore treated as extensions rather than part of the controlled final deliverable.

This paper presents projected ALP sensitivity estimates at the FCC-ee Z pole, focusing on the two experimentally cleanest signatures: the invisible channel (ALP escapes) and

the prompt-resolved channel (ALP decays to two separated photons). The analysis uses a complete simulation chain from matrix-element generation to fast detector response, validated against published experimental results.

Paper structure. Section 2 reviews the relevant ALP physics. Section 3 describes the simulation and detector modelling. Section 4 presents the signal classification, background estimation, and limit-setting method. Section 5 reports the projected exclusion contours. Section 6 discusses the physics reach, limitations, and outlook.

2 Theory

2.1 Analytic model for photophilic ALPs

This analysis assumes a single effective operator, Eq. (1), with all direct ALP couplings to fermions, gluons, W bosons, and Z bosons set to zero at tree level. The only model parameters scanned are therefore the ALP mass m_a and the photon coupling $g_{a\gamma\gamma}$. In this benchmark the total width is taken to be

$$\Gamma_{\text{tot}} \simeq \Gamma(a \rightarrow \gamma\gamma), \quad (2)$$

because $a \rightarrow \gamma\gamma$ is the only open tree-level decay channel in the implemented model. Loop-induced fermion modes are neglected consistently with the photophilic benchmark and with the UFO parameter choice used in the simulation.

2.2 ALP–photon vertex and associated production

The collider event is easiest to read in the same order as it happens. First, the incoming electron and positron annihilate through the ordinary QED vertex into a virtual photon,

$$e^+e^- \rightarrow \gamma^*(q), \quad (3)$$

where $q = p_{e^-} + p_{e^+}$ is the four-momentum carried by the virtual photon. This photon is *off shell*: a real photon must satisfy $q^2 = 0$, but here $q^2 = s$, the squared centre-of-mass energy of the collision. The virtual photon then uses the ALP–photon interaction to produce a real recoil photon and an ALP,

$$\gamma^*(q) \rightarrow \gamma(k) a(p_a). \quad (4)$$

The full associated-production process is therefore

$$e^+e^- \rightarrow \gamma a, \quad (5)$$

with the real photon providing the mono-energetic recoil tag. Photon-fusion production $e^+e^- \rightarrow e^+e^-a$ can have a larger inclusive rate, but it places the scattered electrons in the forward beampipe region and is much harder to separate from radiative Bhabha and beam-induced backgrounds. The analytic and Monte Carlo pipeline therefore targets the recoil-photon topology.

The same ALP-photon vertex controls both this production process and the later ALP decay. Using

$$\tilde{F}^{\mu\nu} = \frac{1}{2}\epsilon^{\mu\nu\rho\sigma}F_{\rho\sigma}, \quad \epsilon^{0123} = +1, \quad (6)$$

the interaction term gives the two-photon vertex

$$V^{\mu\nu}(k_1, k_2) = -i g_{a\gamma\gamma} \epsilon^{\mu\nu\alpha\beta} k_{1\alpha} k_{2\beta}, \quad (7)$$

where both photon momenta are taken incoming to the vertex. Here $F_{\mu\nu}$ is the electromagnetic field-strength tensor, while $\tilde{F}^{\mu\nu}$ is its dual. The object $\epsilon^{\mu\nu\rho\sigma}$ is the four-dimensional Levi-Civita symbol: it is completely antisymmetric, equals +1 for the ordered indices (0, 1, 2, 3), changes sign whenever two indices are exchanged, and vanishes whenever two indices are the same. Its appearance encodes the pseudoscalar, CP-odd nature of the ALP-photon interaction. The Greek indices μ, ν, α, β are Lorentz indices, and repeated indices are summed over according to the Einstein convention. The vectors k_1 and k_2 are the four-momenta of the two photon legs attached to the vertex. In production one of those photon legs is the off-shell $\gamma^*(q)$ and the other is the real recoil photon $\gamma(k)$.

The Born-level amplitude for Eq. (5) is

$$\mathcal{M} = \frac{e g_{a\gamma\gamma}}{s} [\bar{v}(p_{e^+}) \gamma^\mu u(p_{e^-})] \epsilon_{\mu\nu\alpha\beta} q^\alpha k^\beta \varepsilon^{*\nu}(k). \quad (8)$$

The factors in this expression have direct diagrammatic meanings. The electric charge e comes from the ordinary QED vertex where the electron-positron pair annihilates into a virtual photon. The spinors $u(p_{e^-})$ and $\bar{v}(p_{e^+})$ describe the incoming electron and positron, and γ^μ is a Dirac matrix. The factor $1/s$ is the propagator of the intermediate virtual photon after contracting with the photon-current indices. The factor $g_{a\gamma\gamma} \epsilon_{\mu\nu\alpha\beta} q^\alpha k^\beta$ is the ALP-photon vertex from Eq. (7); q is the virtual-photon momentum and k is the outgoing real-photon momentum. Finally, $\varepsilon^{*\nu}(k)$ is the polarization vector of the observed outgoing photon.

2.3 Two-body production kinematics and cross section

In the centre-of-mass frame the total four-momentum is $P^\mu = p_{e^-}^\mu + p_{e^+}^\mu = (\sqrt{s}, 0, 0, 0)$. Momentum conservation gives $P^\mu = p_a^\mu + k^\mu$, with $k^2 = 0$ for the recoil photon and

119 $p_a^2 = m_a^2$ for the ALP. Squaring $P - k = p_a$ gives

$$(P - k)^2 = m_a^2 \quad \Rightarrow \quad s - 2P \cdot k = m_a^2. \quad (9)$$

120 Here p_{e^-} and p_{e^+} are the incoming electron and positron four-momenta, p_a is the outgoing
 121 ALP four-momentum, and k is the outgoing recoil-photon four-momentum. The invariant
 122 $s = P^2 = (p_{e^-} + p_{e^+})^2$ is the squared centre-of-mass energy, so $\sqrt{s}$ is the total collision
 123 energy. Since $P \cdot k = \sqrt{s} E_\gamma$ in the centre-of-mass frame, the recoil photon energy is fixed:

$$E_{\gamma, \text{recoil}} = \frac{s - m_a^2}{2\sqrt{s}}. \quad (10)$$

124 The ALP energy and common three-momentum magnitude are therefore

$$E_a = \sqrt{s} - E_\gamma = \frac{s + m_a^2}{2\sqrt{s}}, \quad |\mathbf{p}_a| = |\mathbf{k}| = \frac{s - m_a^2}{2\sqrt{s}}. \quad (11)$$

125 These relations are the source of the mono-energetic recoil-photon signature. They are
 126 purely kinematic: at fixed $\sqrt{s}$ and m_a , a two-body final state has no freedom to choose
 127 a different recoil-photon energy before detector smearing or initial-state radiation. They
 128 also give the ALP boost,

$$\gamma_a = \frac{E_a}{m_a}, \quad \beta_a \gamma_a = \frac{|\mathbf{p}_a|}{m_a}. \quad (12)$$

129 Averaging over the two electron and two positron spin states and summing over the
 130 final photon polarisation gives, in the limit $m_e^2 \ll s$,

$$|\overline{\mathcal{M}}|^2 = \frac{e^2 g_{a\gamma\gamma}^2}{4} (1 + \cos^2 \theta_{\text{CM}}) \left(1 - \frac{m_a^2}{s}\right)^2, \quad (13)$$

131 where θ_{CM} is the recoil-photon polar angle relative to the incoming electron direction.
 132 The spin average appears because the FCC-ee beams are treated as unpolarized: the
 133 calculation averages over the four possible initial spin configurations and sums over the
 134 unobserved final photon polarization. The factor $(1 + \cos^2 \theta_{\text{CM}})$ gives the angular shape,
 135 while $(1 - m_a^2/s)^2$ reflects the momentum dependence of the pseudoscalar vertex.

136 The two-body phase-space formula for unpolarised $2 \rightarrow 2$ scattering is

$$\frac{d\sigma}{d\Omega} = \frac{1}{64\pi^2 s} \frac{|\mathbf{p}_f|}{|\mathbf{p}_i|} |\overline{\mathcal{M}}|^2, \quad (14)$$

137 with $|\mathbf{p}_i| = \sqrt{s}/2$ and $|\mathbf{p}_f| = (s - m_a^2)/(2\sqrt{s})$. In this expression $d\Omega$ is the solid-angle
 138 element of the outgoing photon, $|\mathbf{p}_i|$ is the magnitude of the incoming electron or positron
 139 three-momentum in the centre-of-mass frame, and $|\mathbf{p}_f|$ is the magnitude of either final-
 140 state momentum. The ratio $|\mathbf{p}_f|/|\mathbf{p}_i|$ is the two-body phase-space suppression: if m_a

approaches $\sqrt{s}$, there is little momentum left for the final state. Integrating over the azimuthal angle then gives

$$\frac{d\sigma}{d\cos\theta_{\text{CM}}} = \frac{\alpha g_{a\gamma\gamma}^2}{32} (1 + \cos^2\theta_{\text{CM}}) \left(1 - \frac{m_a^2}{s}\right)^3. \quad (15)$$

The third power of $(1 - m_a^2/s)$ has a useful interpretation: one power comes from the available two-body phase space and two powers come from the pseudoscalar momentum dependence of the $a\gamma\gamma$ vertex. Finally,

$$\int_{-1}^1 (1 + \cos^2\theta) d\cos\theta = \frac{8}{3}, \quad (16)$$

so the total cross section over the full angular range is

$$\sigma(e^+e^- \rightarrow \gamma a) = \frac{\alpha g_{a\gamma\gamma}^2}{12} \left(1 - \frac{m_a^2}{s}\right)^3. \quad (17)$$

At $\sqrt{s} = 91.2 \text{ GeV}$ and $g_{a\gamma\gamma} = 10^{-3} \text{ GeV}^{-1}$, Eq. (17) gives $\sigma \approx 5.5 \times 10^{-5} \text{ pb}$; at $g_{a\gamma\gamma} = 10^{-6} \text{ GeV}^{-1}$ it falls to $5.5 \times 10^{-11} \text{ pb}$. The enormous FCC-ee Z-pole luminosity ($150 \text{ ab}^{-1} = 1.5 \times 10^8 \text{ pb}^{-1}$) compensates for small couplings and is what makes sensitivity below 10^{-6} GeV^{-1} possible.

2.4 ALP decay to two photons

After the ALP is produced, it can decay through the same $a\gamma\gamma$ vertex into two photons,

$$a(p) \rightarrow \gamma(k_1)\gamma(k_2). \quad (18)$$

The decay amplitude is

$$\mathcal{M} = -g_{a\gamma\gamma} \epsilon^{\mu\nu\alpha\beta} k_{1\alpha} k_{2\beta} \varepsilon_{1\mu}^* \varepsilon_{2\nu}^*. \quad (19)$$

In this equation p is the ALP four-momentum, while k_1 and k_2 are the four-momenta of the two decay photons. $\varepsilon_{1\mu}$ and $\varepsilon_{2\nu}$ are the photon polarization four-vectors. The star means complex conjugation, as appropriate for outgoing external particles. These polarization vectors describe the two physical transverse polarization states of a photon. They should not be confused with the Levi-Civita symbol $\epsilon^{\mu\nu\alpha\beta}$; the different symbol ε^μ is used here to keep the two concepts separate.

The detector does not measure the photon polarization in this analysis. Therefore the observable decay probability must include all allowed final polarization states. This is why we compute $\sum_{\text{pol}} |\mathcal{M}|^2$ rather than a single fixed-polarization amplitude. Because the amplitude is gauge invariant, terms proportional to a photon momentum vanish by the Ward identity, so the physical polarization sums can be replaced by the compact

165 covariant expression $\sum_{\lambda} \varepsilon_{\mu}^{(\lambda)*} \varepsilon_{\mu'}^{(\lambda)} \rightarrow -g_{\mu\mu'}$. With the mostly-minus metric, this gives

$$\begin{aligned} \sum_{\text{pol}} |\mathcal{M}|^2 &= g_{a\gamma\gamma}^2 \epsilon^{\mu\nu\alpha\beta} \epsilon_{\mu\nu}{}^{\alpha'\beta'} k_{1\alpha} k_{2\beta} k_{1\alpha'} k_{2\beta'} \\ &= -2g_{a\gamma\gamma}^2 [k_1^2 k_2^2 - (k_1 \cdot k_2)^2]. \end{aligned} \quad (20)$$

166 The first line follows from multiplying Eq. (19) by its complex conjugate and summing over
167 the two photon polarizations. The second line uses the standard Levi-Civita contraction
168 identity

$$\epsilon^{\mu\nu\alpha\beta} \epsilon_{\mu\nu}{}^{\alpha'\beta'} = -2 \left(g^{\alpha\alpha'} g^{\beta\beta'} - g^{\alpha\beta'} g^{\beta\alpha'} \right). \quad (21)$$

169 The photons are on shell and massless, so $k_1^2 = k_2^2 = 0$. Momentum conservation gives
170 $p = k_1 + k_2$, and because the ALP is on shell, $p^2 = m_a^2$. Therefore $m_a^2 = (k_1 + k_2)^2 = 2k_1 \cdot k_2$,
171 so

$$\sum_{\text{pol}} |\mathcal{M}|^2 = 2g_{a\gamma\gamma}^2 (k_1 \cdot k_2)^2 = \frac{g_{a\gamma\gamma}^2 m_a^4}{2}. \quad (22)$$

172 For correctly paired photons from the ALP decay,

$$M_{\gamma\gamma}^2 = (k_1 + k_2)^2 = 2E_1 E_2 (1 - \cos \theta_{12}) = m_a^2. \quad (23)$$

173 Here E_1 and E_2 are the measured photon energies and θ_{12} is the opening angle between
174 the two photons. This invariant mass is independent of the lab-frame boost of the ALP,
175 so correctly paired decay photons form a peak at $M_{\gamma\gamma} = m_a$. That peak is the analysis
176 observable in the prompt-resolved channel.

177 For a two-body decay into identical massless particles,

$$\Gamma = \frac{1}{2m_a} \frac{1}{2!} \int d\Phi_2 \sum_{\text{pol}} |\mathcal{M}|^2, \quad \int d\Phi_2 = \frac{1}{8\pi}. \quad (24)$$

178 This formula is Fermi's golden rule written in relativistic form. The factor $1/(2m_a)$ is
179 the flux normalization for one decaying particle of mass m_a . The factor $1/2!$ is essential
180 because the two final photons are identical: without it, the same physical final state would
181 be counted twice when the labels 1 and 2 are exchanged. The object $d\Phi_2$ is the Lorentz-
182 invariant two-body phase-space element; for two massless particles integrated over all
183 directions it equals $1/(8\pi)$. Equivalently, $\Gamma = (1/2!) |\mathbf{k}| |\mathcal{M}|^2 / (8\pi m_a^2)$ with $|\mathbf{k}| = m_a/2$,
184 where $|\mathbf{k}|$ is the magnitude of either photon three-momentum in the ALP rest frame.
185 Substituting Eq. (22) gives the width used throughout the project:

$$\Gamma(a \rightarrow \gamma\gamma) = \frac{g_{a\gamma\gamma}^2 m_a^3}{64\pi}. \quad (25)$$

The proper lifetime and proper decay length are then

$$\tau_a = \frac{1}{\Gamma_a}, \quad c\tau_a = \frac{\hbar c}{\Gamma_a}, \quad (26)$$

with $\hbar c = 1.973269804 \times 10^{-16}$ GeV m. The lifetime τ_a is the mean time before decay in the ALP rest frame. Multiplying by c converts this time into the mean proper distance travelled in that rest frame. In the lab frame the decay length is increased by the boost,

$$\ell_a = \beta_a \gamma_a c\tau_a = \frac{|\mathbf{p}_a|}{m_a} \frac{\hbar c}{\Gamma(a \rightarrow \gamma\gamma)}. \quad (27)$$

where $\beta_a = v_a/c$, $\gamma_a = E_a/m_a$, and $\beta_a \gamma_a = |\mathbf{p}_a|/m_a$. This is the quantity that determines whether the ALP decays before the tracker, inside the detector, or outside the calorimeter. Equations (25) and (27) show that ℓ_a decreases with increasing $g_{a\gamma\gamma}$ and increasing m_a : heavy, strongly coupled ALPs decay promptly, while light, weakly coupled ALPs are long-lived and escape the detector.

2.5 Decay probabilities and photon separation

The probability density for an ALP to decay after travelling a lab-frame distance L is exponential,

$$f(L) = \frac{1}{\ell_a} e^{-L/\ell_a}. \quad (28)$$

This is the usual decay law written in distance rather than time. The parameter ℓ_a is the mean lab-frame decay distance from Eq. (27); a larger ℓ_a means that the ALP is more likely to survive to the calorimeter or escape the detector. Integrating this density over the relevant detector length intervals gives the probabilities used in the analytic scan:

$$P(L > L_{\max}) = e^{-L_{\max}/\ell_a}, \quad (29)$$

$$P(L < L_{\min}) = 1 - e^{-L_{\min}/\ell_a}, \quad (30)$$

$$P(L_{\min} < L < L_{\max}) = e^{-L_{\min}/\ell_a} - e^{-L_{\max}/\ell_a}. \quad (31)$$

The first line is the survival probability to the outer detector radius $L_{\max}$ and defines the invisible signature. The second is the probability to decay before the prompt-distance scale $L_{\min}$. The third is the probability to decay between those two scales, producing a displaced signature. These probabilities define the invisible, prompt, and displaced regions of the grid.

The angular resolution boundary follows from boosting the isotropic $a \rightarrow \gamma\gamma$ decay. For an ALP moving along the z axis, the two photons are emitted back-to-back in the ALP rest frame. The smallest lab-frame opening angle occurs when the rest-frame photons are

emitted transverse to the boost, for which

$$\Delta\theta_{\gamma\gamma}^{\min} = 2 \arctan\left(\frac{1}{\gamma_a\beta_a}\right) = 2 \arctan\left(\frac{m_a}{|\mathbf{p}_a|}\right). \quad (32)$$

The physical reason is simple: a highly boosted ALP carries the two decay photons forward in almost the same direction, so the two electromagnetic showers can merge in the calorimeter. Equation (32) uses the exact boost factor through $\gamma_a\beta_a = |\mathbf{p}_a|/m_a$. For light ALPs, $|\mathbf{p}_a| \simeq \sqrt{s}/2$ and $\gamma_a \simeq \sqrt{s}/(2m_a)$, so

$$\Delta\theta_{\gamma\gamma}^{\min} \simeq \frac{2}{\gamma_a} \simeq \frac{4m_a}{\sqrt{s}}. \quad (33)$$

The prompt-resolved channel requires this opening angle to be larger than the detector angular resolution, $\Delta\theta_{\text{res}}$. If $\Delta\theta_{\gamma\gamma}^{\min} < \Delta\theta_{\text{res}}$, the two photons are treated as merged even if the ALP decays promptly.

2.6 Signal signatures

The ALP phenomenology at the detector level is determined by two quantities: the decay length ℓ_a and the minimum diphoton opening angle $\Delta\theta_{\min}$ of the ALP decay products. The opening-angle criterion uses the boosted two-body result derived in Eq. (33). We classify each point $(m_a, g_{a\gamma\gamma})$ in the parameter space into four exclusive signatures:

- **Invisible** ($\ell_a > L_{\text{max}}$): the ALP escapes before decaying. The signature is a single photon with missing energy, $e^+e^- \rightarrow \gamma + \cancel{E}$.
- **Prompt-resolved** ($\ell_a < L_{\min}$ and $\Delta\theta_{\min} \geq \Delta\theta_{\text{res}}$): the ALP decays immediately to two resolvable photons, $e^+e^- \rightarrow \gamma (\rightarrow \gamma\gamma)$. The three-photon final state carries an invariant-mass peak at m_a .
- **Displaced-resolved** ($L_{\min} \leq \ell_a \leq L_{\text{max}}$): the ALP decays inside the detector volume but away from the interaction point.
- **Merged** ($\ell_a \leq L_{\text{max}}$ and $\Delta\theta_{\min} < \Delta\theta_{\text{res}}$): the ALP decay photons are not angularly resolved; the signature resembles a single broad photon cluster.

The invisible and prompt-resolved signatures are the primary targets of this analysis. Displaced and merged signatures require dedicated reconstruction and background models, so they are classified but not included as exclusion contours in the present study.

3 Simulation and Detector Modelling

3.1 Signal generation and computational strategy

Signal events are generated using MADGRAPH5_AMC@NLO [13] with a UFO model implementing the ALP effective field theory of Eq. (1). The dense exclusion scan is not produced by running a full detector simulation at every point in the 180×180 mass–coupling grid. Instead, the code uses a two-level strategy:

1. The full $(m_a, g_{a\gamma\gamma})$ grid is evaluated analytically using Eqs. (17)–(33). This gives the production cross section, lifetime, decay-region probability, recoil energy, and opening angle for each point.
2. A finite set of contour-relevant points is passed through the full detector-level chain. These points are used to validate the matrix element, lifetime convention, Delphes output, and analysis observables, and to build branch-aware detector correction factors.

This separation is possible because associated production contains one ALP photon vertex, so at tree level $\sigma(e^+e^- \rightarrow \gamma a) \propto g_{a\gamma\gamma}^2$ exactly. The shape of the production process is therefore controlled mainly by m_a and $\sqrt{s}$, while the coupling rescales the yield and changes the lifetime.

3.2 Choice of scan ranges

The final FCC-ee projection is evaluated on a logarithmic 180×180 grid,

$$m_a \in [10^{-2}, 80] \text{ GeV}, \quad g_{a\gamma\gamma} \in [10^{-8}, 10^{-1}] \text{ GeV}^{-1}. \quad (34)$$

The lower mass bound, 10 MeV, is chosen because below this value the FCC-ee Z-pole diphoton opening angle $\Delta\theta_{\min} \simeq 4m_a/\sqrt{s}$ is far below the conservative IDEA angular resolution used here. For example, at $m_a = 10 \text{ MeV}$ the opening angle is only 0.025° , compared with $\Delta\theta_{\text{res}} = 1.5^\circ$. Any prompt decay would therefore be a merged-photon object, while weakly coupled ALPs are invisible. Dedicated merged-cluster and very-low-mass beam-dump-style treatments are outside the scope of the current collider projection.

The upper mass bound, 80 GeV, is set by the kinematic and practical reach of associated production at $\sqrt{s} = 91.2 \text{ GeV}$. The hard kinematic condition is $m_a < \sqrt{s}$, but the recoil photon energy becomes small and the $(1 - m_a^2/s)^3$ phase-space factor suppresses the cross section near the endpoint. Stopping at 80 GeV retains the high-mass collider region where the prompt-resolved channel is meaningful, while avoiding a threshold-dominated tail where a full experimental analysis would need a different low-recoil selection.

The coupling interval in Eq. (34) is deliberately wider than the final visible contours. Values below 10^{-8} GeV^{-1} are far below the 150 ab^{-1} event-yield threshold because $\sigma \propto g_{a\gamma\gamma}^2$. Values up to 10^{-1} GeV^{-1} are included to ensure that the invisible upper lifetime boundary is bracketed for the lightest masses. This upper end should not be interpreted as an equally reliable EFT region for every ALP model. In the UFO parameter-card writer we keep $f_a = 1000 \text{ GeV}$ fixed and vary $K_B + K_W$ to realise the requested physical $g_{a\gamma\gamma}$; very large $g_{a\gamma\gamma}$ therefore corresponds to large dimensionless Wilson coefficients and is used mainly as a numerical bracket for the lifetime transition.

The detector-level chain for one signal point is:

$$\begin{aligned} \text{param_card} &\rightarrow \text{MADGRAPH (LHE)} \rightarrow \text{PYTHIA (HepMC)} \\ &\rightarrow \text{DELPHES (ROOT)} \rightarrow \text{Python validation.} \end{aligned} \quad (35)$$

The imported UFO does not expose $g_{a\gamma\gamma}$ directly. The code writes a point-specific parameter card by mapping the physical coupling to the UFO parameters f_a , K_B , and K_W through

$$g_{a\gamma\gamma} = \frac{\alpha_{\text{em}}(K_B + K_W)}{\sqrt{2}\pi f_a}. \quad (36)$$

The UFO also contains an $a\gamma Z$ vertex. In the imported model this vertex is proportional to

$$C_{\gamma Z} \propto K_W c_w^2 - K_B s_w^2, \quad (37)$$

where s_w and c_w are computed from the same electroweak inputs as the UFO. The default parameter-card writer therefore uses

$$K_B = (K_B + K_W) c_w^2, \quad K_W = (K_B + K_W) s_w^2, \quad (38)$$

which makes Eq. (37) vanish at tree level while keeping the $a\gamma\gamma$ coupling of Eq. (36). This choice is essential for the interpretation of the project. It isolates the one-parameter photophilic benchmark, makes the MadGraph production normalisation comparable to Eq. (17) and to the Belle II closure, and prevents the Z-pole projection from being contaminated by resonant $Z \rightarrow \gamma a$ production. Including a genuine ALP–Z coupling is a separate physics question with a different operator, different rates, and different signal interpretation.

For lifetime-dependent results, the code writes the project convention

$$\Gamma(a \rightarrow \gamma\gamma) = \frac{g_{a\gamma\gamma}^2 m_a^3}{64\pi} \quad (39)$$

directly into the `DECAY 9999` block of the parameter card. Pythia then decays the ALP to $\gamma\gamma$ with the corresponding proper lifetime. The resulting HepMC event record is passed to DELPHES 3 [15] for fast detector simulation using the IDEA detector card [10].

3.3 Repository implementation and data products

The repository mirrors the physics pipeline. Analytic formulae and validation gates live in the theory layer; MadGraph, Pythia, HepMC, and Delphes wrappers live in the MC layer; background building, signal-efficiency correction, projection solving, and plotting live in the analysis layer. The main machine-readable inputs are the locked JSON files specifying the Belle II closure and FCC-ee Z-pole assumptions. The main checked-in outputs are:

- the analytic theory grid;
- the Belle II closure curve, plot, and JSON summary;
- binned FCC-ee background histograms for the resolved and invisible channels;
- the collected 284-point detector-level ALP signal summary;
- the branch-aware full-analysis efficiency map;
- the FCC-ee projection CSV and signature-classification grid;
- the full, close-up, and combined ALP-photon money plots.

Large raw production files such as LHE, HepMC, ROOT, and full Condor work directories are deliberately not versioned. They are reproducible intermediate products, while the compact CSV, JSON, PNG, and PDF outputs are the report-facing deliverables.

3.4 Detector modelling

The IDEA detector concept [10] is a proposed FCC-ee detector featuring a solenoidal magnetic field, silicon tracking, a drift chamber, and a crystal electromagnetic calorimeter. The key detector parameters used in this analysis are:

- Photon acceptance: $|\eta| < 3.0$, $E_\gamma > 0.5 \text{ GeV}$;
- Photon reconstruction efficiency: 99%;
- Diphoton angular resolution: $\Delta\theta_{\text{res}} = 1.5^\circ$ (conservative calorimeter-cell scale, $\pi/120$);
- Tracker inner radius (L_{min}): 0.02 m;
- Tracker/calorimeter outer radius (L_{max}): 2.5 m.

These values are locked in `analysis/configs/fccee_zpole_inputs.json`. The final deliverable uses the assumptions in Table 1; changing any of these inputs requires regenerating the classification grid, contours, and money plots.

3.5 Validation gates

Before applying the analysis to FCC-ee kinematics, the pipeline is checked by a set of validation gates. These gates are run by the repository validation scripts and by the full signal-production wrapper. They test both the physics formulae and the file handoff between programs.

The Belle II closure is the most important external validation anchor. Belle II operates at $\sqrt{s} = 10.58 \text{ GeV}$ with 445 pb^{-1} and uses the same process $e^+e^- \rightarrow \gamma a$, $a \rightarrow \gamma\gamma$. The procedure infers the effective signal-yield threshold implied by the digitised Belle II boundary from AxionLimits [12]. It then checks that the project production cross section, decay width, lifetime convention, and prompt/resolved region logic reproduce the same contour. The closure test yields

$$\max_{m_a} \left| \log_{10} \left(\frac{g_{\text{closure}}}{g_{\text{published}}} \right) \right| = 7.6 \times 10^{-3}, \quad (40)$$

well within the 2% tolerance used as the acceptance criterion. This demonstrates that the analytic model, lifetime convention, and signal-region logic are all internally consistent and agree with a published experimental result.

3.6 SM background samples

Background events are simulated at $\sqrt{s} = 91.2 \text{ GeV}$ with MADGRAPH5_AMC@NLO + PYTHIA 8 + DELPHES/IDEA. Two channels are generated:

- **Resolved-prompt background:** $e^+e^- \rightarrow \gamma\gamma\gamma$ (SM QED triple-photon production), with $\sigma = 7.306 \text{ pb}$ and $N_{\text{gen}} = 10\,000$ events.
- **Invisible background:** $e^+e^- \rightarrow \gamma\nu\bar{\nu}$ (radiative neutrino production), with $\sigma = 134.9 \text{ pb}$ and $N_{\text{gen}} = 10\,000$ events.

Backgrounds are normalised to $\mathcal{L} = 150 \text{ ab}^{-1}$. At full luminosity, the expected number of background events across the analysis histograms is 2.58×10^9 (resolved) and 5.43×10^9 (invisible). The resolved background is the dominant QED irreducible process for the three-photon topology; the invisible background is the leading electroweak process with a photon plus missing energy.

4 Analysis

4.1 Signature classification grid

Each of the $180 \times 180 = 32\,400$ grid points $(m_a, g_{a\gamma\gamma})$ is classified into one of the four signatures using the analytic formulae of Section 2. With the IDEA detector assumptions,

the grid decomposes as follows (Table 3):

The prompt-resolved region dominates at high $g_{a\gamma\gamma}$ and intermediate m_a , where the ALP decays promptly and the photons are resolvable. The invisible region covers most of the low-coupling parameter space, where the decay length exceeds the detector radius. The prompt-resolved threshold mass is set by $\Delta\theta_{\min} = \Delta\theta_{\text{res}}$:

$$m_a^{\text{thresh}} = \frac{\sqrt{s} \Delta\theta_{\text{res}}}{4} \approx \frac{91.2 \text{ GeV} \times 0.0262}{4} \approx 0.597 \text{ GeV}. \quad (41)$$

Below this mass, the photon pair from ALP decay would be reconstructed as a single cluster, falling into the merged category.

4.2 Signal efficiency and detector corrections

The expected number of signal events for a given grid point is:

$$N_{\text{sig}} = \mathcal{L} \cdot \sigma(m_a, g_{a\gamma\gamma}) \cdot P_{\text{region}}(m_a, g_{a\gamma\gamma}) \cdot \epsilon_{\text{param}} \cdot C_{\text{Delphes}}, \quad (42)$$

where P_{region} is the analytic probability that the ALP falls in the target signature region:

$$P_{\text{invisible}} = e^{-L_{\text{max}}/\ell_a}, \quad (43)$$

$$P_{\text{prompt}} = 1 - e^{-L_{\text{min}}/\ell_a}. \quad (44)$$

The parametric efficiency ϵ_{param} accounts for angular acceptance and photon reconstruction. For the invisible channel it is the single-photon acceptance times one photon-efficiency factor; for the prompt-resolved channel it is the same acceptance times three photon-efficiency factors. The remaining detector response is absorbed into the Delphes-derived correction factor. The detector correction factor C_{Delphes} is derived from a channel-aware scan of 284 signal points through the full DELPHES simulation. For each point, the full analysis pipeline—including the signal observable histogram, background comparison, and analysis bin acceptance—is compared to the analytic expectation. The resulting correction is interpolated as a function of ALP mass separately for each branch. Across the signal scan, the mean detector correction factors are:

- Invisible lower branch: $\langle C \rangle = 0.998$ (range 0.969–1.003),
- Invisible upper branch: $\langle C \rangle = 7.8 \times 10^6$ (extremely large correction from the non-monotonic g dependence of the invisible yield; this branch is numerically fragile),
- Prompt-resolved: $\langle C \rangle = 1.02$ (range 0.900–2.62).

4.3 Limit-setting method

A projected 90% confidence level (CL) exclusion is derived using a binned Asimov $\Delta\chi^2$ test [16]. For a given mass hypothesis, the analysis observable is:

- Resolved channel: the diphoton invariant mass $M_{\gamma\gamma}$, with the signal modelled as a Gaussian centred at m_a with width $\sigma_M = \max(0.05 m_a, 0.05 \text{ GeV})$;
- Invisible channel: the recoil photon energy E_γ , with the signal modelled as a Gaussian centred at $(s - m_a^2)/(2\sqrt{s})$ with width $\sigma_E = \max(0.05 E_\gamma, 0.5 \text{ GeV})$. The histogram uses 264 bins over $[0, 50] \text{ GeV}$, extending beyond the nominal $\sqrt{s}/2 = 45.6 \text{ GeV}$ endpoint to retain detector- and ISR-smeared signal events.

For each mass point, the minimum signal yield required for exclusion is:

$$N_{\text{sig}}^{\text{req}} = \max \left(3, \sqrt{\frac{2.71}{\sum_i f_i^2 / \max(B_i, 1)}} \right), \quad (45)$$

where f_i is the normalised signal fraction in bin i , B_i is the expected SM background in that bin, and 3 is a conservative event-count floor. The threshold $\Delta\chi^2 = 2.71$ corresponds to a one-sided 90% CL. The contour in the $(m_a, g_{a\gamma\gamma})$ plane is then the locus of points where $N_{\text{sig}}(m_a, g_{a\gamma\gamma}) = N_{\text{sig}}^{\text{req}}(m_a)$.

Because the invisible signal yield is non-monotonic in $g_{a\gamma\gamma}$ —it peaks at intermediate coupling and falls off both towards small coupling (fewer events produced) and large coupling (ALP decays before reaching L_{max})—the invisible contour has two branches: the *lower branch* (sensitivity floor from statistics) and the *upper branch* (sensitivity ceiling from lifetime).

5 Results

5.1 Signature map

Figure 1 shows the classification of the $(m_a, g_{a\gamma\gamma})$ plane under the IDEA detector assumptions. The four signatures occupy well-separated regions. The invisible region (blue) dominates at low coupling over a wide mass range; the prompt-resolved region (red) covers the high-coupling, intermediate-mass space above the resolved threshold; the merged region (orange) occupies the low-mass, high-coupling corner where photon pairs are not angularly resolved; and the displaced-resolved band (green) forms a narrow transition strip. The physics boundaries between regions—the $\ell_a = L_{\text{min}}$ and $\ell_a = L_{\text{max}}$ curves computed from Eqs. (25) and (27)—are overlaid as dashed lines and agree with the classification boundaries to machine precision.




Figure 1: ALP signature classification in the $(m_a, g_{a\gamma\gamma})$ plane for the FCC-ee Z pole ($\sqrt{s} = 91.2 \text{ GeV}$) under IDEA detector assumptions. Dashed lines show the analytic $\ell_a = L_{\min}$ (prompt/displaced boundary) and $\ell_a = L_{\max}$ (displaced/invisible boundary) curves. The thin dash-dot curve shows the angular-resolution boundary $\Delta\theta = \Delta\theta_{\text{res}} = 1.5^\circ$. Only the invisible and prompt-resolved signatures are used in the projected exclusion contours.

5.2 Projected exclusion contours

Figure 2 shows the main result: the projected FCC-ee Z-pole exclusion contours overlaid on the current ALP-photon constraint landscape from AxionLimits [12].

The external regions in Fig. 2 are used as a landscape, not as a single statistically homogeneous exclusion. AxionLimits combines direct laboratory searches, stellar-energy-loss bounds, cosmological arguments, dark matter-motivated projections, and QCD-axion reference bands. These inputs do not all share the same assumptions about ALP abundance, branching fractions, early-universe history, or statistical treatment. The FCC-ee curves overlaid here are different in character: they are direct collider projections for the photophilic benchmark of Eqs. (1) and (38), with no assumption that the ALP is dark matter or that it solves the strong CP problem. Regions labelled by cosmology, stellar physics, and QCD-axion model lines are therefore important context for the global ALP landscape, but they do not replace the controlled collider question answered by this repository.

Invisible channel. The invisible sensitivity spans $m_a \in [0.01, 0.92] \text{ GeV}$. The lower contour (sensitivity floor) sits at:

$$g_{a\gamma\gamma}^{\text{inv}} \approx 5.5 \times 10^{-7} \text{ GeV}^{-1} \quad (\text{nearly flat across the mass range}), \quad (46)$$




Figure 2: ALP-photon parameter space with FCC-ee Z-pole projected sensitivity. *Left*: full ALP-photon landscape from meV to TeV ALP mass. Individual existing constraints are shown as labelled coloured regions (stellar/astrophysical in green, helioscope CAST in pink, colliders and beam experiments in red/blue tones). *Right*: FCC-ee Z-pole zoom region (10 MeV–80 GeV). Shaded blue band: projected invisible exclusion region, bounded below (solid dashed line, sensitivity floor) and above (lighter dash, lifetime ceiling). Shaded green region: projected prompt-resolved exclusion, bounded by the sensitivity contour (green dashed). All constraints shown are existing published bounds. Background constraints from AxionLimits [12].

representing a factor ~ 10 – 100 improvement over the current best collider bounds in this mass range. The near-flatness of the lower contour reflects the dominance of the signal statistics floor (3 events) over the background shape at high luminosity. The upper contour rises from $g_{a\gamma\gamma} \approx 1.3 \times 10^{-6} \text{ GeV}^{-1}$ at $m_a = 0.01 \text{ GeV}$ to $g_{a\gamma\gamma} \approx 0.055 \text{ GeV}^{-1}$ at $m_a = 0.92 \text{ GeV}$, marking the transition to the short-lifetime (prompt) regime.

Prompt-resolved channel. The prompt-resolved sensitivity spans $m_a \in [0.61, 80] \text{ GeV}$. The lower mass threshold is set by Eq. (41). The exclusion contour runs from:

$$g_{a\gamma\gamma}^{\text{res}} \in [1.1 \times 10^{-5}, 2.9 \times 10^{-4}] \text{ GeV}^{-1} \quad (\text{at } m_a = 0.61\text{--}80 \text{ GeV}). \quad (47)$$

In the mass range 1–10 GeV, the prompt-resolved sensitivity at FCC-ee extends below the current Belle II and BaBar limits by roughly one to two orders of magnitude.

Comparison to existing limits. Table 4 summarises the FCC-ee reach at selected masses compared to the best current bounds.

6 Discussion

6.1 Physics reach

The FCC-ee Z-pole presents a qualitatively different sensitivity landscape from any existing experiment. At low ALP masses ($m_a \lesssim 0.1 \text{ GeV}$), the invisible channel exploits the huge luminosity to push the sensitivity floor below $g_{a\gamma\gamma} \sim 10^{-6} \text{ GeV}^{-1}$, a regime otherwise accessible only to stellar observations (HB stars, SN 1987A). At intermediate masses (1–10 GeV), the prompt-resolved channel cleanly covers the region between current B -factory limits and LEP. At high masses (10–80 GeV), the prompt-resolved channel overlaps with and significantly improves on LEP and LHC constraints, benefiting from the clean e^+e^- environment and the mono-energetic photon tag.

The invisible channel’s sensitivity has a fundamental lower limit set by the 3-event floor in Eq. (45): at exactly 3 required events and $\mathcal{L} = 150 \text{ ab}^{-1}$, the cross section threshold corresponds to $\sigma_{\min} = 3/(1.5 \times 10^8 \text{ pb}^{-1}) = 2 \times 10^{-8} \text{ pb}$. From Eq. (17), this maps to $g_{a\gamma\gamma} \approx 5 \times 10^{-7} \text{ GeV}^{-1}$, consistent with Eq. (46). A nearly background-free invisible analysis at this luminosity would be statistics-limited. In the implemented binned-background analysis, the background shape raises the required signal yield above three events for some masses (for example to about 10.7 events at the lightest masses), but the result remains driven mainly by the enormous FCC-ee luminosity and the $\sigma \propto g_{a\gamma\gamma}^2$ scaling.

6.2 Confidence in the projection

The results should be interpreted as a detector-corrected phenomenological projection, not as an experimental limit. Within that scope, the confidence is highest where the analysis is controlled by validated analytic formulae and moderate detector corrections, and lower where small lifetime probabilities are amplified by large correction factors.

The most robust headline statements are therefore the existence of strong FCC-ee reach in the invisible lower branch and in the prompt-resolved branch, the location of the resolved threshold near 0.6 GeV, and the qualitative separation of invisible, prompt/resolved, merged, and displaced regions. The least robust feature is the precise position of the invisible upper branch, which should be described as the expected lifetime boundary rather than as a precision contour.

6.3 Limitations and caveats

Several important caveats apply to the current projections.

Invisible upper branch. The invisible upper contour (the lifetime ceiling) carries very large detector correction factors ($C_{\text{Delphes}} \sim 10^6\text{--}10^8$ in parts of the low-mass tail). These arise because the analytic invisible probability $P_{\text{inv}} = e^{-L_{\text{max}}/\ell_a}$ plunges steeply as $g_{a\gamma\gamma}$

increases, while the Delphes simulation shows a slightly smoother falloff due to ISR and finite detector response. The upper branch should be treated as directional rather than precise.

Detector systematics. The current analysis includes only the dominant SM backgrounds for each channel and does not account for: systematic uncertainties on photon reconstruction efficiency and energy scale; machine-induced backgrounds from beamstrahlung and synchrotron radiation (relevant at FCC-ee); pileup-like machine backgrounds; and higher-order QED and electroweak corrections to the SM background rates.

Merged and displaced signatures. The merged and displaced-resolved signatures together cover 24% of the $(m_a, g_{a\gamma\gamma})$ grid. Merged photons require dedicated cluster-shape reconstruction, and displaced diphotons require dedicated vertex-level analysis. Including these channels would extend the claimed reach to couplings where the ALP is not fully prompt, potentially filling in the gap between the invisible and prompt-resolved contours.

Single-iteration Delphes correction. The branch-aware detector correction factor is derived from signal points on the first-iteration projected contour, not at the corrected contour positions. A second Delphes campaign at the corrected contour points would be needed for a fully self-consistent result.

Belle II closure scope. The closure test validates the analysis model at published-contour level. It does not access Belle II’s private likelihood, background spectra, or efficiency maps, and should not be read as a full reimplementation of the Belle II result.

6.4 Outlook

Several extensions would strengthen and broaden the analysis.

Displaced and merged channels. Implementing dedicated reconstruction for displaced and merged topologies would extend the FCC-ee ALP coverage to a nearly complete $(m_a, g_{a\gamma\gamma})$ map from the resolved threshold ($m_a \approx 0.6 \text{ GeV}$) all the way to the kinematic endpoint ($m_a \lesssim \sqrt{s}/2 = 45.6 \text{ GeV}$), including the intermediate-lifetime region currently left blank.

Higher-order corrections. ISR corrections and NLO QCD/EW effects on the background processes could shift the background normalisation by $\mathcal{O}(10\%)$, which at this level of signal sensitivity may become relevant.

ALP- Z coupling. The present analysis considers only the ALP-photon coupling $g_{a\gamma\gamma}$. ALPs may also couple to the Z boson, modifying both production and decay at FCC-ee energies [5, 6]. In this repository that coupling is cancelled at tree level by Eq. (38). This is why the projection does not use the potentially resonant $Z \rightarrow \gamma a$ enhancement

from the 5×10^{12} Z bosons expected at the FCC-ee Z pole. Including a nonzero $C_{\gamma Z}$ would enrich the phenomenology, but it would be a different two-coupling analysis with additional processes such as $Z \rightarrow \gamma a$ and possibly $e^+e^- \rightarrow Za$.

Other FCC-ee run stages. FCC-ee will also operate at the WW threshold, HZ threshold, and $t\bar{t}$ threshold. The WW and HZ stages provide different kinematics and potentially larger cross sections for ALPs coupling to W bosons or via Higgs loops, while the higher $\sqrt{s}$ extends the prompt-resolved reach to higher m_a .

7 Conclusion

We have computed projected 90% CL exclusion contours for photophilic ALPs at the FCC-ee Z-pole ($\sqrt{s} = 91.2 \text{ GeV}$, $\mathcal{L} = 150 \text{ ab}^{-1}$) using a complete simulation chain from matrix-element generation to fast detector response. For the invisible channel, the analysis achieves sensitivity to $g_{a\gamma\gamma} \approx 5.5 \times 10^{-7} \text{ GeV}^{-1}$ across the mass range $m_a \in [0.01, 0.92] \text{ GeV}$, surpassing current collider bounds by one to two orders of magnitude. For the prompt-resolved channel, the contour covers $m_a \in [0.61, 80] \text{ GeV}$ at $g_{a\gamma\gamma} \sim 10^{-5} - 3 \times 10^{-4} \text{ GeV}^{-1}$, significantly exceeding the reach of existing e^+e^- and LHC experiments.

The analysis is validated against the published Belle II exclusion boundary to 0.8%, confirming the consistency of the analytic production model, lifetime convention, and signal-region logic. The simulation chain uses MadGraph5, Pythia 8, and a Delphes IDEA card; the validation suite also checks the MadGraph cross section, the ALP width and Pythia lifetime, the Delphes output, and the invariant-mass or recoil-energy observable for every detector-level signal point. The most robust parts of the result are the invisible lower branch and the prompt-resolved contour. The invisible upper branch should be interpreted more qualitatively because it is a rapidly varying lifetime boundary with large detector-correction factors.

The results establish FCC-ee as an extraordinarily powerful probe of photophilic ALPs in the 10 MeV–80 GeV mass range, potentially surpassing the world-leading sensitivity of stellar and astrophysical bounds in a controlled collider environment. Future work should extend the analysis to displaced and merged signatures, include machine-induced backgrounds and systematic uncertainties, and incorporate ALP couplings beyond the photon vertex.

Acknowledgements

The author thanks the Nikhef computing cluster for providing the resources used for the full-statistics Monte Carlo campaigns, and the AxionLimits repository [12] for the constraint landscape data.

References

- [1] R. D. Peccei and H. R. Quinn, *CP conservation in the presence of pseudoparticles*, Phys. Rev. Lett. **38**, 1440 (1977).
- [2] S. Weinberg, *A new light boson?*, Phys. Rev. Lett. **40**, 223 (1978).
- [3] F. Wilczek, *Problem of strong P and T invariance in the presence of instantons*, Phys. Rev. Lett. **40**, 279 (1978).
- [4] J. Jaeckel and A. Ringwald, *The low-energy frontier of particle physics*, Annu. Rev. Nucl. Part. Sci. **60**, 405 (2010), [[arXiv:1002.0329](#)].
- [5] I. Brivio *et al.*, *ALPs effective field theory and collider signatures*, Eur. Phys. J. C **77**, 572 (2017), [[arXiv:1701.05379](#)].
- [6] M. J. Dolan, T. Erber, F. Kahlhoefer, and M. Schmidt-Hoberg, *Searching for ALPs with LEP, the LHC and a future circular collider*, JHEP **10**, 079 (2017), [[arXiv:1709.03453](#)].
- [7] J. Jaeckel and M. Thormaehlen, *Distinguishing axion models with the axion-photon coupling*, Phys. Rev. D **95**, 015010 (2017), [[arXiv:1612.00034](#)].
- [8] A. Abada *et al.* (FCC Collaboration), *FCC-ee: The lepton collider*, Eur. Phys. J. Special Topics **228**, 261 (2019), [[arXiv:1905.05078](#)].
- [9] A. Abada *et al.* (FCC Collaboration), *FCC physics opportunities*, Eur. Phys. J. C **79**, 474 (2019), [[arXiv:1905.05078](#)].
- [10] A. Abada *et al.*, *IDEA: A detector concept for FCC-ee*, Eur. Phys. J. Special Topics **228**, 1107 (2019).
- [11] I. Adachi *et al.* (Belle II Collaboration), *Search for axion-like particles produced in e^+e^- collisions at Belle II*, Phys. Rev. Lett. **130**, 231801 (2023), [[arXiv:2206.07598](#)].
- [12] C. O’Hare, *cajohare/AxionLimits: AxionLimits*, Zenodo (2020), DOI: 10.5281/zenodo.3932430, <https://cajohare.github.io/AxionLimits/>.
- [13] J. Alwall *et al.*, *The automated computation of tree-level and next-to-leading order differential cross sections, and their matching to parton showers*, JHEP **07**, 079 (2014), [[arXiv:1405.0301](#)].
- [14] T. Sjöstrand *et al.*, *Introduction to PYTHIA 8.2*, Comput. Phys. Commun. **191**, 159 (2015), [[arXiv:1410.3012](#)].

[15] J. de Favereau *et al.* (DELPHES 3 Collaboration), *DELPHES 3: A modular framework for fast simulation of a generic detector*, JHEP **02**, 057 (2014), [[arXiv:1307.6346](#)].

[16] G. Cowan, K. Cranmer, E. Gross, and O. Vitells, *Asymptotic formulae for likelihood-based tests of new physics*, Eur. Phys. J. C **71**, 1554 (2011), [[arXiv:1007.1727](#)].

Appendix: Contribution Statement

Iñaki Pardo led the physics study, repository organization, simulation campaigns, analysis development, validation, plotting, and writing. The project was developed in collaboration with project partners and with Codex-assisted implementation and documentation work. External software, models, detector cards, constraint data, and papers are credited in the repository provenance report and bibliography.

Table 1: Locked assumptions used for the final FCC-ee Z-pole deliverable.

Ingredient	Value	Role in the analysis
Collider stage	$\sqrt{s} = 91.2 \text{ GeV}$, $\mathcal{L} = 150 \text{ ab}^{-1}$	Baseline FCC-ee Z-pole scenario; sets recoil energy, production phase space, and event normalisation.
Physics benchmark	Photophilic ALP, $C_{\gamma Z} = 0$ at tree level	Keeps the result a one-coupling $g_{a\gamma\gamma}$ projection rather than a mixed photon-Z benchmark.
Scan domain	$m_a = 0.01\text{--}80 \text{ GeV}$, $g_{a\gamma\gamma} = 10^{-8}\text{--}10^{-1} \text{ GeV}^{-1}$	Covers the FCC-ee mass reach and brackets both the lower sensitivity floor and upper invisible lifetime boundary.
Signal regions	Invisible and prompt-resolved only	These are the two channels with simple robust observables: recoil-photon energy and $M_{\gamma\gamma}$.
Detector lengths	$L_{\min} = 0.02 \text{ m}$, $L_{\max} = 2.5 \text{ m}$	Define prompt, displaced, and invisible decay probabilities.
Photon reconstruction	$ \eta < 3.0$, $E_\gamma > 0.5 \text{ GeV}$, $\epsilon_\gamma = 0.99$	Implements the IDEA-like photon acceptance and reconstruction efficiency.
Angular resolution	$\Delta\theta_{\text{res}} = 1.5^\circ$	Defines whether $a \rightarrow \gamma\gamma$ is resolved or merged.
Analysis resolutions	$\sigma_M = \max(0.05m_a, 0.05 \text{ GeV})$, $\sigma_E = \max(0.05E_\gamma, 0.5 \text{ GeV})$	Widths of the Gaussian signal templates in the resolved and invisible binned likelihoods.
SM backgrounds	$e^+e^- \rightarrow \gamma\gamma\gamma$ and $e^+e^- \rightarrow \gamma\nu\bar{\nu}$	Irreducible backgrounds for the prompt-resolved and invisible channels, respectively.
Statistical criterion	One-sided 90% CL, $\Delta\chi^2 = 2.71$, with a 3-event floor	Converts binned signal-plus-background expectations into exclusion contours.
External constraints	AxionLimits commit 7d375f4879b32406 a239fe48d2615a4b fd9bc0bb	Provides the published ALP, QCD-axion, astrophysical, cosmological, and laboratory landscape used only as context for the money plots.

Table 2: Validation gates implemented in the analysis pipeline.

Gate	What is tested	Current status
Software smoke test	A generic $e^+e^- \rightarrow \mu^+\mu^-$ chain proves that MadGraph, Pythia, HepMC, Delphes, ROOT, and the Python reader work together.	Integrated as the cluster setup test.
Gate 1: cross section	The MadGraph integrated weight for $e^+e^- \rightarrow \gamma a$ is compared to Eq. (17) at the generated $(m_a, g_{a\gamma\gamma}, \sqrt{s})$ point.	Passed for 284/284 detector-level ALP scan points.
Gate 2: width and lifetime	The UFO width convention is checked, the project 64π width is written into the parameter card, and the Pythia input lifetime is checked against Eq. (27).	Convention locked to Eq. (25); Pythia summaries are checked per point.
Gate 3: Belle II closure	The Belle II public contour is reconstructed using the same production, lifetime, and prompt/resolved logic as the FCC-ee projection.	Passed with maximum log residual 7.6×10^{-3} .
Detector-observable gate	Resolved points must reconstruct a diphoton mass near m_a ; invisible points must reconstruct the recoil photon near Eq. (10).	Passed for 284/284 detector-level ALP scan points.
Background gate	SM background ROOT outputs are converted to binned luminosity-normalised histograms used directly by the limit solver.	Two full-stat 10 000-event background channels included.

 Table 3: Signature classification of the $(m_a, g_{a\gamma\gamma})$ grid.

Signature	Grid points	Fraction
Prompt-resolved	14 171	43.7%
Invisible	10 452	32.3%
Merged	5 989	18.5%
Displaced-resolved	1 788	5.5%
Total	32 400	100%

Table 4: FCC-ee projected sensitivity versus current best limits at representative masses. Current limits from AxionLimits [12].

m_a	Channel	Current best $g_{a\gamma\gamma}$ [GeV ⁻¹]	FCC-ee projected $g_{a\gamma\gamma}$ [GeV ⁻¹]
0.1 GeV	Invisible	$\sim 10^{-4}$ (beam dumps)	5.5×10^{-7}
1 GeV	Prompt-resolved	$\sim 10^{-3}$ (Belle II)	$\sim 3 \times 10^{-5}$
10 GeV	Prompt-resolved	$\sim 10^{-4}$ (BaBar)	$\sim 2 \times 10^{-5}$
50 GeV	Prompt-resolved	$\sim 10^{-3}$ (LEP)	$\sim 2 \times 10^{-4}$

Table 5: Reliability assessment of the main ingredients in the projection.

Ingredient	Confidence	Reason
Production normalisation	High	The associated-production cross section is analytic and Gate 1 agrees with MadGraph for all 284 detector-level signal points.
Width and lifetime convention	High	The 64π width convention is locked in the parameter-card writer, Pythia receives the same lifetime, and Gate 2 checks the convention.
Belle II closure logic	High at public-contour level	Gate 3 reproduces the digitised Belle II boundary to a maximum residual of 7.6×10^{-3} in $\log_{10} g_{a\gamma\gamma}$.
Prompt-resolved FCC-ee contour	Moderate to high	The branch has reasonable Delphes correction factors (median close to one, range 0.90–2.62) and a direct invariant-mass validation at detector level.
Invisible lower contour	Moderate to high	The lower branch is controlled by production rate, survival probability near unity, recoil-energy validation, and correction factors close to one.
Invisible upper contour	Low to moderate	This is a lifetime ceiling where the invisible probability changes extremely rapidly with $g_{a\gamma\gamma}$; the detector-correction factors become very large in the low-mass tail.
Absolute experimental reach	Moderate	The study includes leading SM backgrounds and Delphes response, but not full detector systematics, machine backgrounds, or a real FCC-ee experimental likelihood.